

# Scientific Machine Learning for the No-Three-In-Line Problem: Energy-Based Continuous Dynamics and Neural ODEs

## 1 Motivation and Role in the Overall Framework

Previous work in this project has compared three main paradigms for the No-Three-In-Line problem:

- **Integer Linear Programming (ILP):** Provides provably optimal solutions but scales poorly with grid size.
- **PatternBoost Transformer:** Learns patterns from data and achieves near-optimal performance on moderate grids.
- **Reinforcement Learning (PPO):** Explores the combinatorial space via sequential decisions; effective on small grids but struggles as constraints tighten.

The SciML extension adds a fourth viewpoint:

*Treat the discrete grid as a continuous field evolving under an ODE / Neural ODE, with a carefully designed energy functional whose minima correspond to good (or optimal) configurations.*

This gives:

- A physics-inspired, interpretable formulation.
- A continuous relaxation with polynomial-time ODE integration.
- A natural basis for hybrid methods (e.g., using SciML solutions as warm starts for ILP or RL).

## 2 Continuous Relaxation and State Representation

We relax the binary grid to continuous variables.

### 2.1 Grid and variables

Let  $n$  be the grid size. The discrete decision variables are:

$$x_{ij} \in \{0, 1\}, \quad i, j = 1, \dots, n.$$

We introduce a continuous relaxation:

$$x_{ij}(t) \in [0, 1], \quad i, j = 1, \dots, n, \quad t \in [0, T],$$

and collect them into a vector:

$$\mathbf{x}(t) \in [0, 1]^{n^2}.$$

Interpretation:

- $x_{ij}(t) \approx 1$ : cell  $(i, j)$  contains a point.
- $x_{ij}(t) \approx 0$ : cell  $(i, j)$  is empty.

Our goal is to define dynamics

$$\dot{\mathbf{x}}(t) = f_\theta(\mathbf{x}(t), t)$$

that drive  $\mathbf{x}(t)$  toward low-energy configurations that are dense, nearly binary, and obey the “no three in a line” constraint.

### 3 Energy Functional Design

Let  $L$  denote the set of all triplets of collinear grid positions:

$$L = \left\{ ((i_1, j_1), (i_2, j_2), (i_3, j_3)) \mid \text{three distinct collinear cells} \right\}.$$

The energy  $E(\mathbf{x})$  combines three terms.

#### 3.1 Collinearity penalty

We penalize any triple of cells that are simultaneously active:

$$E_{\text{collinear}}(\mathbf{x}) = \sum_{((i_1, j_1), (i_2, j_2), (i_3, j_3)) \in L} \alpha x_{i_1 j_1} x_{i_2 j_2} x_{i_3 j_3}, \quad (1)$$

where  $\alpha > 0$  controls the strength of the penalty. If three cells along a line are close to 1, their product becomes large and the energy increases.

#### 3.2 Point-count reward

We reward the presence of points by decreasing the energy when  $x_{ij}$  is large:

$$E_{\text{count}}(\mathbf{x}) = -\beta \sum_{i=1}^n \sum_{j=1}^n x_{ij}, \quad (2)$$

with  $\beta > 0$ . Minimizing this term pushes the system toward grids with many active cells.

#### 3.3 Binary regularization

To avoid trivial fractional configurations (e.g.,  $x_{ij} \approx 0.5$  everywhere), we use a double-well potential:

$$E_{\text{binary}}(\mathbf{x}) = \gamma \sum_{i=1}^n \sum_{j=1}^n x_{ij}^2 (1 - x_{ij})^2, \quad (3)$$

with  $\gamma > 0$ . This term has minima at  $x_{ij} = 0$  and  $x_{ij} = 1$ , encouraging near-binary states.

### 3.4 Total energy

The total energy is:

$$E(\mathbf{x}; \alpha, \beta, \gamma) = E_{\text{collinear}}(\mathbf{x}) + E_{\text{count}}(\mathbf{x}) + E_{\text{binary}}(\mathbf{x}). \quad (4)$$

The scalars  $\alpha, \beta, \gamma$  can be:

- hand-tuned for a simple baseline, or
- learned via SciML as trainable parameters.

## 4 Dynamics: Gradient Flow and Neural ODE

### 4.1 Pure gradient-flow ODE

A natural choice is gradient flow in continuous time:

$$\dot{\mathbf{x}}(t) = -\nabla_{\mathbf{x}} E(\mathbf{x}(t)). \quad (5)$$

This is continuous-time gradient descent on the energy landscape. Fixed points correspond to local minima of  $E$ .

In SciML, we implement this as an ODE:

$$\frac{d}{dt} \mathbf{x}(t) = f(\mathbf{x}(t), t; \alpha, \beta, \gamma),$$

where  $f$  computes  $-\nabla E$  using automatic differentiation.

### 4.2 Physics-informed Neural ODE

To increase expressive power, we can add a learned correction:

$$\dot{\mathbf{x}}(t) = -\nabla_{\mathbf{x}} E(\mathbf{x}(t)) + g_{\theta}(\mathbf{x}(t), t), \quad (6)$$

where  $g_{\theta}$  is a neural network (e.g., MLP or convolutional network over the grid), parameterized by  $\theta$ .

This yields a *physics-informed Neural ODE*, where:

- $-\nabla E$  encodes known structure and constraints.
- $g_{\theta}$  learns to improve convergence and escape poor local minima.

## 5 Training Objectives

We describe two training regimes, one fully unsupervised and one teacher-guided using ILP solutions.

## 5.1 Unsupervised energy-minimization training

In unsupervised mode, we simply ask the system to learn parameters that produce low-energy configurations from random initial states.

Let  $\mathbf{x}^{(b)}(t)$  denote the solution trajectory for the  $b$ -th initial condition, and let  $\mathbf{x}^{(b)}(T)$  be the final state at terminal time  $T$ . The unsupervised loss is:

$$\mathcal{L}_{\text{energy}}(\theta, \alpha, \beta, \gamma) = \frac{1}{B} \sum_{b=1}^B E(\mathbf{x}^{(b)}(T); \alpha, \beta, \gamma), \quad (7)$$

where  $B$  is the batch size. Parameters are updated by backpropagating through the ODE using adjoint sensitivity methods.

## 5.2 Teacher-guided training with ILP or PatternBoost

When optimal (or near-optimal) discrete solutions  $\hat{\mathbf{y}}^{(k)} \in \{0, 1\}^{n^2}$  are available from ILP or PatternBoost for small grids, we can add a teacher loss.

For the  $k$ -th training instance, with final state  $\mathbf{x}^{(k)}(T)$ , define:

$$\mathcal{L}_{\text{teacher}} = \frac{1}{K} \sum_{k=1}^K \|\sigma(\mathbf{x}^{(k)}(T)) - \hat{\mathbf{y}}^{(k)}\|_2^2, \quad (8)$$

where  $\sigma(\cdot)$  is a smooth squashing function (e.g., elementwise sigmoid), and  $K$  is the number of teacher-labeled instances in a batch.

The total loss becomes:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{energy}} + \lambda_{\text{teacher}} \mathcal{L}_{\text{teacher}}, \quad (9)$$

with a hyperparameter  $\lambda_{\text{teacher}} \geq 0$  controlling the strength of teacher guidance.

## 6 From Continuous to Discrete Solutions

After solving the ODE from  $t = 0$  to  $t = T$ , we obtain  $\mathbf{x}(T)$ . To obtain a discrete grid:

### 1. Thresholding:

$$\hat{x}_{ij} = \begin{cases} 1, & x_{ij}(T) \geq \tau, \\ 0, & x_{ij}(T) < \tau, \end{cases}$$

with threshold  $\tau \in (0, 1)$  (e.g.,  $\tau = 0.5$ ).

### 2. Local repair (optional):

For any line that contains three or more active points, we greedily deactivate points to satisfy the constraint, prioritizing removals that minimize the decrease in total point count.

This yields a feasible configuration that can be evaluated against the ILP optimum.

## 7 Training-Time and Complexity Analysis

Let:

- $n$ : grid size.

- $d = n^2$ : number of continuous variables.
- $|L|$ : number of collinear triples; asymptotically  $|L| = O(n^3)$ .
- $N_t$ : number of time steps (RHS evaluations) used by the ODE solver per trajectory.
- $B$ : batch size (number of trajectories per iteration).
- $N_{\text{iter}}$ : number of training iterations.

## 7.1 Cost per ODE solve

Each evaluation of the vector field requires computing the gradient  $\nabla_{\mathbf{x}} E(\mathbf{x})$ .

- Computing  $E_{\text{collinear}}(\mathbf{x})$  is  $O(|L|)$  since we sum over all collinear triples.
- With automatic differentiation, computing  $\nabla_{\mathbf{x}} E(\mathbf{x})$  is also  $O(|L|)$  up to a constant factor.

The ODE solver requires  $N_t$  evaluations, so:

$$\text{Cost per ODE solve} \approx O(N_t |L|) = O(N_t n^3).$$

## 7.2 Cost per training iteration and total cost

With batch size  $B$ , a single training iteration requires  $B$  ODE solves:

$$\text{Cost per iteration} \approx O(B N_t n^3).$$

Over  $N_{\text{iter}}$  iterations:

$$\text{Total training cost} \approx O(N_{\text{iter}} B N_t n^3).$$

## 7.3 Qualitative comparison

- **ILP**: Worst-case exponential in problem size; empirically hits practical limits around  $n = 19$  for exact solutions.
- **PatternBoost / PPO**: Training cost grows with number of episodes, model size, and horizon length; the combinatorial action space becomes challenging at large  $n$ .
- **SciML (this work)**: Training dominated by ODE solves with cost  $O(n^3)$  per trajectory, plus backpropagation. This offers a different scaling profile and potentially better behavior than ILP for moderate-to-large grids.

This analysis can be directly reported in the ICML paper to highlight the trade-offs between methods.

## 8 Experimental Plan for SciML

For each grid size  $n \in \{8, 10, 12, 14, 16, 19\}$ , we propose to evaluate:

- **SciML-GF:** Pure gradient-flow dynamics ( $g_\theta \equiv 0$ ).
- **SciML-NODE:** Physics-informed Neural ODE with non-zero  $g_\theta$ .

For each method and each  $n$ :

1. Run  $R$  trajectories from different random initial conditions (e.g.,  $R = 100$ ).
2. Integrate the ODE to terminal time  $T$ , apply thresholding and optional repair.
3. Measure:
  - Feasible success rate (fraction of runs with zero collinear triples).
  - Average number of points in feasible solutions.
  - Best-of- $R$  number of points.
  - Average training time per iteration and total training time.
  - Inference time per trajectory (single ODE solve).

These metrics can be compared to ILP (optimal where available), PatternBoost, and PPO.

## 9 Pseudocode

### 9.1 Precomputing collinear triples

---

**Algorithm 1** Compute collinear triples  $L$  for an  $n \times n$  grid

---

```
1: function COMPUTECOLLINEARTRIPLIES( $n$ )
2:    $L \leftarrow \emptyset$ 
3:   coords  $\leftarrow \{(i, j) \mid i = 1..n, j = 1..n\}$ 
4:   for each triple of distinct points  $(p_1, p_2, p_3)$  from coords do
5:     if COLLINEAR( $p_1, p_2, p_3$ ) then
6:       Append  $(p_1, p_2, p_3)$  to  $L$ 
7:     end if
8:   end for
9:   return  $L$ 
10: end function
```

---

Two points  $(x_1, y_1)$ ,  $(x_2, y_2)$ ,  $(x_3, y_3)$  are collinear if the triangle area is zero:

$$x_1(y_2 - y_3) + x_2(y_3 - y_1) + x_3(y_1 - y_2) = 0.$$

### 9.2 Energy and vector field

### 9.3 Training loop (high-level)

---

**Algorithm 2** Energy and ODE vector field

---

```
1: function ENERGY( $\mathbf{x}, L, \alpha, \beta, \gamma$ )
2:    $E_{\text{col}} \leftarrow 0$ 
3:   for each triple  $(p_1, p_2, p_3) \in L$  do
4:      $i_1 \leftarrow \text{INDEXOF}(p_1)$ 
5:      $i_2 \leftarrow \text{INDEXOF}(p_2)$ 
6:      $i_3 \leftarrow \text{INDEXOF}(p_3)$ 
7:      $E_{\text{col}} \leftarrow E_{\text{col}} + \alpha x_{i_1} x_{i_2} x_{i_3}$ 
8:   end for
9:    $E_{\text{count}} \leftarrow -\beta \sum_k x_k$ 
10:   $E_{\text{binary}} \leftarrow \gamma \sum_k x_k^2 (1 - x_k)^2$ 
11:  return  $E_{\text{col}} + E_{\text{count}} + E_{\text{binary}}$ 
12: end function
13: function VECTORFIELD( $\mathbf{x}, t, \theta, \alpha, \beta, \gamma, L$ )
14:    $g \leftarrow \nabla_{\mathbf{x}} \text{ENERGY}(\mathbf{x}, L, \alpha, \beta, \gamma)$ 
15:   physics  $\leftarrow -g$ 
16:   if  $\theta$  is used then
17:     correction  $\leftarrow g_{\theta}(\mathbf{x}, t)$ 
18:   else
19:     correction  $\leftarrow 0$ 
20:   end if
21:   return physics + correction
22: end function
```

---

---

**Algorithm 3** SciML Training Loop

---

```
1: Initialize parameters  $\theta, \alpha, \beta, \gamma$ 
2:  $L \leftarrow \text{COMPUTECOLLINEARTRIPLES}(n)$ 
3: for iter = 1 to  $N_{\text{iter}}$  do
4:    $\mathcal{L} \leftarrow 0$ 
5:   for  $b = 1$  to  $B$  do
6:     Sample random initial state  $\mathbf{x}_0^{(b)}$ 
7:     Solve ODE:  $\mathbf{x}^{(b)}(T) \leftarrow \text{SOLVEODE}(\text{VECTORFIELD}, \mathbf{x}_0^{(b)}, [0, T])$ 
8:      $\mathcal{L} \leftarrow \mathcal{L} + \text{ENERGY}(\mathbf{x}^{(b)}(T), L, \alpha, \beta, \gamma)$ 
9:     if teacher data is available then
10:       Sample teacher solution  $\hat{\mathbf{y}}^{(b)}$ 
11:        $\mathcal{L} \leftarrow \mathcal{L} + \lambda_{\text{teacher}} \|\sigma(\mathbf{x}^{(b)}(T)) - \hat{\mathbf{y}}^{(b)}\|_2^2$ 
12:     end if
13:   end for
14:    $\mathcal{L} \leftarrow \mathcal{L}/B$ 
15:   Compute gradients of  $\mathcal{L}$  w.r.t.  $\theta, \alpha, \beta, \gamma$ 
16:   Update parameters with chosen optimizer (e.g., Adam)
17: end for
```

---

## 10 Julia SciML Skeleton (Simulation Only)

Below is a compact, self-contained Julia example that:

- builds an  $n \times n$  grid,
- precomputes collinear triples,
- defines an energy functional,
- implements gradient-flow dynamics via finite-difference gradient (for clarity), and
- solves the ODE from a random initial condition and thresholds the final state.

You can replace the finite-difference gradient with automatic differentiation (e.g., Zygote) and wrap this in DiffEqFlux for full training.

```
using OrdinaryDiffEq
using Random

# -----
# 1. Grid & index helpers
# -----
const n = 6 # example grid size
const N = n * n # total variables

# Map (i, j) -> linear index in 1..N
linear_index(i, j) = (i - 1) * n + j

# -----
# 2. Precompute collinear triples
# -----
function compute_collinear_triples(n::Int)
    triples = Tuple{Tuple{Int,Int}, Tuple{Int,Int}, Tuple{Int,Int}}[]
    coords = [(i, j) for i in 1:n for j in 1:n]

    # Helper: check collinearity of three lattice points
    function collinear(p1, p2, p3)
        (x1, y1) = p1
        (x2, y2) = p2
        (x3, y3) = p3
        # Area of triangle == 0 -> collinear
        return (x1*(y2 - y3) + x2*(y3 - y1) + x3*(y1 - y2)) == 0
    end

    # Enumerate all distinct 3-combinations of points
    L = length(coords)
    for a in 1:(L-2), b in (a+1):(L-1), c in (b+1):L
        p1 = coords[a]
        p2 = coords[b]
        p3 = coords[c]
        if collinear(p1, p2, p3)
            push!(triples, (p1, p2, p3))
        end
    end
end
```

```

        return triples
end

const L_triples = compute_collinear_triples(n)

# -----
# 3. Energy function
# -----
function energy(x::Vector{Float64};
                ::Float64 = 5.0,
                ::Float64 = 1.0,
                ::Float64 = 2.0)

    # Collinearity term
    E_col = 0.0
    @inbounds for (p1, p2, p3) in L_triples
        i1 = linear_index(p1[1], p1[2])
        i2 = linear_index(p2[1], p2[2])
        i3 = linear_index(p3[1], p3[2])
        E_col += * x[i1] * x[i2] * x[i3]
    end

    # Count term (reward more points)
    E_count = - * sum(x)

    # Binary regularization (double well at 0 and 1)
    E_bin = 0.0
    @inbounds for k in 1:length(x)
        E_bin += x[k]^2 * (1.0 - x[k])^2
    end
    E_bin *=

    return E_col + E_count + E_bin
end

# -----
# 4. Approximate gradient via finite differences
# (for demonstration; replace with AD in real experiments)
# -----
function grad_energy_fd(x::Vector{Float64};
                           h::Float64 = 1e-5,
                           ::Float64 = 5.0,
                           ::Float64 = 1.0,
                           ::Float64 = 2.0)

    g = similar(x)
    E0 = energy(x; =, =, =)
    for k in 1:length(x)
        x[k] += h
        Ek = energy(x; =, =, =)
        x[k] -= h
        g[k] = (Ek - E0) / h
    end
    return g

```

```

end

# -----
# 5. ODE vector field: pure gradient flow
# -----
function f!(dx, x, p, t)
    , , = p
    g = grad_energy_fd(x; =, =, =)
    @. dx = -g # gradient descent direction
end

# -----
# 6. Solve the ODE from a random initial condition
# -----
Random.seed!(1234)
x0 = rand(N) # random in (0,1)

tspan = (0.0, 5.0) # time horizon
p = (5.0, 1.0, 2.0) # (, , )
prob = ODEProblem(f!, x0, tspan, p)
sol = solve(prob, Tsit5(); reltol=1e-4, abstol=1e-6)

# Extract final state and threshold
xT = Array(sol(end))
= 0.5
x_binary = [val >= ? 1 : 0 for val in xT]
grid = reshape(x_binary, (n, n))

println("Final discrete grid (1 = point, 0 = empty):")
for i in 1:n
    println(grid[i, :])
end

```

This skeleton can be extended in several directions:

- Replace `grad_energy_fd` with an automatic differentiation-based gradient.
- Wrap the ODE in a `DiffEqFlux` training loop to learn  $\alpha, \beta, \gamma$  or parameters of a neural correction term  $g_\theta$ .
- Scale up  $n$  and integrate this module into your ILP / PatternBoost / PPO comparison.